


# Introduction to performance analysis in openSUSE®

Using Perf



**Tony Jones** – Senior Software Engineer  
SUSE Labs Portland, Oregon USA [tonyj@suse.com](mailto:tonyj@suse.com)

# Give yourself a break

- Performance analysis is not easy
- It's rarely black and white, more shades of grey.
- Lots of subjectivity



# Methodologies and practice are useful

- Methodologies are useful, there are lots of them.
- → be methodical
- If you have to dig in, know what options exist.
- Write code to test hypothesis/understand tools.



# There are lots of tools, perf is one

- And it's not the first one you want to use.
  - dmesg/syslog
  - top
  - iostat/vmstat
  - ltop
  - pidstat
  - strace



# Performance Counters

- Hardware resource to aid performance analysis.
- Available on x86 since Pentium 3.
- Availability via CPUID and MSR
- Basic PMU functionality “standardized” with arch\_perfmon\_v1 thru v4
- Beyond this, additional PMUs micro-architecture specific



# Perfmon

- Perfmon 2.0 designed by HP for ia64.
- Still in kernel tree for ia64.
- Extended for other architectures as Perfmon 2.X.
- Submission to LKML in 2008 led to counter submission of Performance Counters For Linux
- libpfm v4 still relevant today



# Perf

- Introduced in 2009 in as Performance Counters for Linux, kernel version 2.6.31
- Now known as ‘perf events’
- A swiss-army knife but still mostly focused on cpu usage and integration with tracing.



# Where to begin

- Install perf :)
- perf cmd
  - top
  - list
  - record
  - report
  - annotate
  - trace





# Available events

- perf list

# perf list	
cache-misses	[Hardware event]
cpu-cycles OR cycles	[Hardware event]
instructions	[Hardware event]
...	
context-switches OR cs	[Software event]
cpu-migrations OR migrations	[Software event]
...	
L1-dcache-loads	[Hardware cache event]
...	
dTLB-load-misses	[Hardware cache event]
...	
rNNN	[Raw hardware event descriptor]
...	
ext4:ext4_alloc_da_blocks	[Tracepoint event]
...	
module:module_free	[Tracepoint event]
...	
net:netif_rx	[Tracepoint event]
...	
sched:sched_switch	[Tracepoint event]
...	
signal:signal_deliver	[Tracepoint event]
..	
syscalls:sys_enter_chroot	[Tracepoint event]
..	
syscalls:sys_exit_kexec_load	[Tracepoint event]

# Event monikers

```
intel# grep "model name" /proc/cpuinfo | uniq
model name      : Intel(R) Xeon(R) CPU E5-2420 v2 @ 2.20GHz
```

```
intel# ls /sys/devices/cpu/events/
branch-instructions  cache-misses      instructions  ref-cycles
branch-misses        cache-references  mem-loads    stalled-cycles-frontend
bus-cycles           cpu-cycles        mem-stores
```

```
intel# cat /sys/devices/cpu/events/cpu-cycles
event=0x3c
```

```
# /usr/bin/showevtinfo > /tmp/events
```

```
IDX          : 37748736
PMU name     : ix86arch (Intel X86 architectural PMU)
Name         : UNHALTED_CORE_CYCLES
Equiv        : None
Flags        : None
Desc         : count core clock cycles whenever the clock signal on the specific core is running (not
halted)
Code         : 0x3c
Modif-00     : 0x00 : PMU : [k] : monitor at priv level 0 (boolean)
Modif-01     : 0x01 : PMU : [u] : monitor at priv level 1, 2, 3 (boolean)
Modif-02     : 0x02 : PMU : [e] : edge level (may require counter-mask >= 1) (boolean)
Modif-03     : 0x03 : PMU : [i] : invert (boolean)
Modif-04     : 0x04 : PMU : [c] : counter-mask in range [0-255] (integer)
Modif-05     : 0x05 : PMU : [t] : measure any thread (boolean)
```

# Event monikers (cont)

```
amd# grep "model name" /proc/cpuinfo | uniq
model name      : AMD Opteron(tm) Processor 6128
```

```
amd# ls /sys/devices/cpu/events/
branch-instructions  cache-misses      cpu-cycles
branch-misses        cache-references  instructions
```

```
amd# cat /sys/devices/cpu/events/cpu-cycles
event=0x76
```

```
# /usr/bin/showevtinfo > /tmp/events
```

```
[amd]
IDX      : 134217766
PMU name : amd64_fam10h_istanbul (AMD64 Fam10h Istanbul)
Name     : CPU_CLK_UNHALTED
Equiv    : None
Flags    : None
Desc     : CPU Clocks not Halted
Code     : 0x76
Modif-00 : 0x00 : PMU : [k] : monitor at priv level 0 (boolean)
Modif-01 : 0x01 : PMU : [u] : monitor at priv level 1, 2, 3 (boolean)
Modif-02 : 0x02 : PMU : [e] : edge level (boolean)
Modif-03 : 0x03 : PMU : [i] : invert (boolean)
Modif-04 : 0x04 : PMU : [c] : counter-mask in range [0-255] (integer)
Modif-05 : 0x05 : PMU : [h] : monitor in hypervisor (boolean)
Modif-06 : 0x06 : PMU : [g] : measure in guest (boolean)
```



# Sampling vs counting

- Counters operate in two distinct ways
  - Counted.
  - Sampled.
- Sampling rate can be specified via either
  - Frequency (-F)
  - Count (-c)



# Perf top

- Perf top. Sampling view similar to standard UNIX utility
- Provides view into where cpu is spending time. Not everything is user cpu bound (0.1% user, 4.1% sys)

Samples: 32K of event 'cycles:ppp', Event count (approx.): 8956122825

Overhead	Shared Object	Symbol
44.93%	[kernel]	[k] __raw_spin_lock
36.76%	[kernel]	[k] sync_inodes_sb
1.78%	[kernel]	[k] __raw_spin_unlock
0.78%	[kernel]	[k] __filemap_fdatawait_range
0.58%	[kernel]	[k] find_get_pages_tag
0.56%	[kernel]	[k] _atomic_dec_and_lock
0.50%	[kernel]	[k] __iget
0.45%	[kernel]	[k] nmi
0.41%	[kernel]	[k] iput
0.36%	[kernel]	[k] unmap_page_range
0.32%	[kernel]	[k] __wake_up_bit
0.25%	[kernel]	[k] filemap_map_pages
0.23%	[kernel]	[k] clear_page_c_e
0.22%	[kernel]	[k] copy_page

...  
...

# Perf top (cont)

- Sometimes it is (top reports 100% user, 0% sys)

```
Samples: 867K of event 'cycles:ppp', Event count (approx.): 340340242278
Overhead Shared Object      Symbol
 51.02% libcrypto.so.1.0.0  [.] DES_set_key_unchecked
 26.95% libcrypto.so.1.0.0  [.] DES_encrypt1
  1.07% libcrypto.so.1.0.0  [.] DES_set_odd_parity
  0.53% libcrypto.so.1.0.0  [.] EVP_DigestInit_ex
  0.44% libc-2.23.so        [.] __int_free
  0.39% libcrypto.so.1.0.0  [.] EVP_Digest
  0.38% libc-2.23.so        [.] __int_malloc
  0.30% libcrypto.so.1.0.0  [.] EVP_MD_CTX_cleanup
  0.29% libc-2.23.so        [.] malloc
...
```

# Perf top (cont)

- Possible to annotate hot code
- Enable callgraphs (-g) to show paths to hot code
- Two types of callgraphs for userspace. Frame pointer and dwarf

```
# zcat /proc/config.gz | egrep -e "CONFIG_FRAME_POINTER|UNWIND"  
CONFIG_UNWIND_INFO=y  
CONFIG_STACK_UNWIND=y  
# CONFIG_FRAME_POINTER is not set
```



# Perf record

- Samples counters for later analysis.
- Lots of command line options:
  - perf record *command*
  - perf record -a *command*
  - perf record -a
- Default scope is kernel+usermode
  - perf record -e cpu-cycles:u *command*
  - perf record -a -e cpu-cycles:k





# Perf record

```
# perf record -c 1000 -e instructions:u ./linpack 200
```

```
# perf report -D | grep -c PERF_RECORD_SAMPLE  
3048011
```

```
# perf report -n --stdio
```

```
# Total Lost Samples: 0
```

```
#
```

```
# Samples: 3M of event 'instructions:u'
```

```
# Event count (approx.): 3048011000
```

```
#
```

#	Overhead	Samples	Command	Shared Object	Symbol
#	.....	.....	.....	.....	.....
	48.04%	1464144	linpack	linpack	[.] daxpy_r
	40.31%	1228752	linpack	linpack	[.] daxpy_ur
	7.72%	235359	linpack	linpack	[.] matgen
	2.05%	62548	linpack	linpack	[.] dgefa
	0.83%	25291	linpack	linpack	[.] idamax
	0.52%	15798	linpack	linpack	[.] dscal_r
	0.41%	12496	linpack	linpack	[.] dscal_ur
	0.11%	3337	linpack	linpack	[.] dgesl
	0.00%	56	linpack	libc-2.23.so	[.] _dl_addr



# Perf record

```
# perf annotate --print-line --stdio daxpy_r
...
      :          /* code for both increments equal to 1 */
      :
      :      for (i = 0; i < n; i++)
0.21 :      40200e:      movl    $0x0, -0x4(%rbp)
0.00 :      402015:      jmp     40206e <daxpy_r+0x163>
      :          dy[i] = dy[i] + da*dx[i];
linpack.c:599 10.24 :      402017:      mov     -0x4(%rbp),%eax
      :      40201a:      cltq
      :      40201c:      lea     0x0(,%rax,8),%rdx
0.01 :      402024:      mov     -0x30(%rbp),%rax
0.30 :      402028:      add     %rdx,%rax
10.31 :      40202b:      mov     -0x4(%rbp),%edx
      :      40202e:      movslq  %edx,%rdx
0.04 :      402031:      lea     0x0(,%rdx,8),%rcx
0.33 :      402039:      mov     -0x30(%rbp),%rdx
10.37 :      40203d:      add     %rcx,%rdx
      :      402040:      movsd   (%rdx),%xmm1
linpack.c:599 3.04 :      402044:      mov     -0x4(%rbp),%edx
linpack.c:599 9.35 :      402047:      movslq  %edx,%rdx
      :      40204a:      lea     0x0(,%rdx,8),%rcx
0.00 :      402052:      mov     -0x28(%rbp),%rdx
      :      402056:      add     %rcx,%rdx
linpack.c:599 1.28 :      402059:      movsd   (%rdx),%xmm0
linpack.c:599 9.58 :      40205d:      mulsd   -0x20(%rbp),%xmm0
linpack.c:599 0.88 :      402062:      addsd   %xmm1,%xmm0
linpack.c:599 10.86 :      402066:      movsd   %xmm0, (%rax)
      :
17.50 :
```



# Perf record

```
# perf record -c 1000 -e instructions ./linpack 200
# perf report -n --stdio | head -30
```

```
# Total Lost Samples: 0
```

```
#
```

```
# Samples: 1M of event 'instructions'
```

```
# Event count (approx.): 1868359000
```

```
#
```

#	Overhead	Samples	Command	Shared Object	Symbol
	47.46%	886656	linpack	linpack	[.] daxpy_r
	40.03%	747911	linpack	linpack	[.] daxpy_ur
	8.32%	155453	linpack	linpack	[.] matgen
	1.60%	29808	linpack	linpack	[.] dgefa
	0.79%	14756	linpack	linpack	[.] idamax
	0.49%	9101	linpack	linpack	[.] dscal_r
	0.39%	7298	linpack	linpack	[.] dscal_ur
	0.26%	4825	linpack	[kernel.kallsyms]	[k] _raw_spin_lock_irq
	0.22%	4181	linpack	[kernel.kallsyms]	[k] __local_bh_enable
	0.10%	1833	linpack	[kernel.kallsyms]	[k] run_timer_softirq
	0.08%	1537	linpack	linpack	[.] dgesl
	0.03%	610	linpack	[kernel.kallsyms]	[k] __softirqentry_text_start
	0.03%	599	linpack	[kernel.kallsyms]	[k] idle_cpu
	0.01%	275	linpack	[kernel.kallsyms]	[k] load_balance
	0.01%	253	linpack	[kernel.kallsyms]	[k] _raw_spin_unlock_irq
	0.01%	226	linpack	[kernel.kallsyms]	[k] update_sd_lb_stats
	0.01%	159	linpack	[kernel.kallsyms]	[k] _raw_spin_unlock_irqrestore
	0.01%	159	linpack	[kernel.kallsyms]	[k] select_task_rq_fair
	0.01%	157	linpack	[kernel.kallsyms]	[k] enqueue_entity



# What events should I monitor

- Instructions per cycle

```
# perf stat -e instructions:u -e cycles:u ./linpack 200
```

```
Performance counter stats for './linpack 200':
```

```
77,693,957,078      instructions:u          #    2.73  insn per cycle
```

```
28,421,726,540      cycles:u
```

```
25.760008514 seconds time elapsed
```



# What events should I monitor

- Cache misses relative to references

```
#define MULT 8 // 1 == 8 * 1024 * 1024, ~8MB of 15MB L3 cache
int main()
{
    int *p, i; const int num = MULT * 1024 * 1024, loops=50;
```

```
    if (!(p=malloc(num * sizeof(int))))
        return 1;
```

```
    srand((unsigned)time(NULL));
```

```
    for (i;i<num*loops;i++)
        p[rand() % num] += rand();
```

```
    free(p);
    return 0;
}
```

```
# perf stat -e cache-misses:u -e cache-references:u ./badcache
Performance counter stats for './badcache':
```

```
227,020,317      cache-misses:u      # 54.474 % of all cache refs
416,750,100      cache-references:u
```

```
23.962244212 seconds time elapsed
```



# What events should I monitor

- Cache misses relative to references

```
IDX      : 216006703
PMU name : ivb_ep (Intel Ivy Bridge EP)
Name     : MEM_LOAD_UOPS_RETIRED
Flags    : [precise]
Desc     : Memory loads uops retired
Code     : 0xd1
Umask-00 : 0x40 : PMU : [HIT_LFB] : [precise] : A load missed L1D but hit the Fill Buffer
Umask-01 : 0x01 : PMU : [L1_HIT] : [precise] : Load hit in nearest-level (L1D) cache
Umask-02 : 0x02 : PMU : [L2_HIT] : [precise] : Load hit in mid-level (L2) cache
Umask-03 : 0x10 : PMU : [L2_MISS] : [precise] : Load misses in mid-level (L2) cache
Umask-04 : 0x04 : PMU : [L3_HIT] : [precise] : Load hit in last-level (L3) cache with no snoop needed
Umask-05 : 0x20 : PMU : [L3_MISS] : [precise] : Load miss in last-level (L3) cache

# for i in L2_HIT L2_MISS L3_HIT L3_MISS; do echo -n "`/usr/bin/evt2raw MEM_LOAD_UOPS_RETIRED:$i` "; done
r5302d1 r5310d1 r5304d1 r5320d1

# perf stat -e r5302d1:u -e r5310d1:u -e r5304d1:u -e r5320d1:u ./badcache

Performance counter stats for './badcache':

      2,875,282      r5302d1:u      (50.00%)
    412,616,740      r5310d1:u      (50.01%)
    187,720,675      r5304d1:u      (50.00%)
    224,533,511      r5320d1:u      (49.99%)

36.058470947 seconds time elapsed
```

# What events should I monitor

- Branch misses relative to branch instructions

```
enum{size = 8}; int t[size] = {0,0,0,0,0,0,0,0};

void func(int _v)
{
#define DO(x) if (_v == x) t[x]++
#define EDO(x) else DO(x)

    DO(0); EDO(1); EDO(2); EDO(3); EDO(4); EDO(5); EDO(6); EDO(7);
}

int main()
{
    int i;
    volatile int v;
    const int num = 20 * 1024 * 1024;

    srand((unsigned)time(NULL));

    for (i=0;i<num;i++){
        v=rand() % size;
        func(v);
    }

    //for (i=0;i<size;i++)
    //    printf("t[%d]=%d\n", i, t[i]);
}
```



# What events should I monitor

- Counted

```
# perf stat -e branches:u -e branch-misses:u branch_8x
```

```
Performance counter stats for 'branch_8x':
```

```
490,265,094      branches:u
19,435,491       branch-misses:u      #    3.96% of all branches

1.361357178 seconds time elapsed
```





# What events should I monitor

- Sampled

```
# perf record -e branches:u -e branch-misses:u branch_8x

# perf report -n --stdio --dsos branch_8x
# dso: branch_8x
#
# Samples: 2K of event 'branches:u'
# Event count (approx.): 490409129
# Overhead      Samples  Command    Symbol
   42.50%         1196  branch_8x  [.] func
   21.67%          609  branch_8x  [.] main

# Samples: 2K of event 'branch-misses:u'
# Event count (approx.): 19453325
# Overhead      Samples  Command    Symbol
   73.30%         2058  branch_8x  [.] func
   23.79%          669  branch_8x  [.] main
```



# Sampling skew/skid

- When a performance counter overflows, IP passed.
- On heavily pipelined architectures reported IP may not be correct
- PEBS, precise event based sampling
- Specified via the precision qualifier (0 to 3)
  - event:p [1, skid shall be constant]
  - event:pp [2, request 0 skid]
  - event:ppp [3, require 0 skid]



# Tracing and tracepoints

- perf list tracepoint

```
# perf list tracepoint
List of pre-defined events (to be used in -e):
  block:block_bio_backmerge          [Tracepoint event]
  ...
  sched:sched_switch                 [Tracepoint event]
  ...
  syscalls:sys_enter_accept          [Tracepoint event]
  ...
  syscalls:sys_exit_accept           [Tracepoint event]

# cat /sys/kernel/debug/tracing/events/sched/sched_switch/format
name: sched_switch
ID: 273
format:
  field:unsigned short common_type;   offset:0;   size:2;   signed:0;
  field:unsigned char common_flags;   offset:2;   size:1;   signed:0;
  ...
print fmt: "prev_comm=%s prev_pid=%d prev_prio=%d prev_state=%s ==> next_comm=%s
next_pid=%d next_prio=%d", REC->prev_comm, REC->prev_pid, REC->prev_prio, REC->prev_state
& (2048-1) ? __print_flags(REC->prev_state & (2048-1), "|", { 1, "S" }, { 2, "D" }, { 4,
"T" }, { 8, "t" }, { 16, "Z" }, { 32, "X" }, { 64, "x" }, { 128, "K" }, { 256, "W" },
{ 512, "P" }, { 1024, "N" }) : "R", REC->prev_state & 2048 ? "+" : "", REC->next_comm,
REC->next_pid, REC->next_prio
```



# Tracing and tracepoints

- `perf stat -a -e 'syscalls:sys_enter_*' sleep 30`
- `perf stat -e 'sched:sched_switch' ./benchmark`
- `perf record -a -g -e 'syscalls:sys_enter_write' --filter 'count > 1024'`

```
# tail -1 /sys/kernel/debug/tracing/events/syscalls/sys_enter_write/format
print fmt: "fd: 0x%08lx, buf: 0x%08lx, count: 0x%08lx", ((unsigned long)(REC->fd)),
((unsigned long)(REC->buf)), ((unsigned long)(REC->count))
```



# Perf trace

- Track costly syscalls

```
# perf trace --duration 1.0 dd if=/dev/zero of=/tmp/out bs=1k count=10
10+0 records in
10+0 records out
10240 btes (10 kB, 10 KiB) copied, 0.000448134 s, 22.9 MB/s
```

```
# perf trace --duration 1.0 dd if=/dev/zero of=/tmp/out bs=1k count=10 conv=fdatasync
8.652 ( 6.035 ms): dd/9820 fdatasync(fd: 1</tmp/out>
10+0 records in
10+0 records out
10240 bytes (10 kB, 10 KiB) copied, 0.00663256 s, 1.5 MB/s
```

```
# perf trace --duration 1.0 dd if=/dev/zero of=/tmp/out bs=1k count=10 oflag=sync
7.795 ( 6.127 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
9.278 ( 1.463 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
10.746 ( 1.448 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
12.195 ( 1.418 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
13.704 ( 1.486 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
15.312 ( 1.564 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
16.754 ( 1.419 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
18.214 ( 1.437 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
19.741 ( 1.503 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
21.200 ( 1.431 ms): dd/9847 write(fd: 1</tmp/out>, buf: 0x559fbf625000, count: 1024 ) = 1024
```



# Perf trace

- Summary of syscalls

```
# perf trace -s dd if=/dev/zero of=/tmp/out bs=1k count=10
10+0 records in
10+0 records out
10240 bytes (10 kB, 10 KiB) copied, 0.000364407 s, 28.1 MB/s
```

Summary of events:

dd (9990), 305 events, 97.8%, 0.000 msec

syscall	calls	total (msec)	min (msec)	avg (msec)	max (msec)	stddev (%)
read	13	0.032	0.002	0.002	0.004	9.14%
write	13	0.102	0.004	0.008	0.021	17.20%
open	50	0.257	0.003	0.005	0.057	21.01%
close	22	0.037	0.001	0.002	0.002	2.77%
fstat	17	0.037	0.002	0.002	0.005	9.72%
lseek	1	0.002	0.002	0.002	0.002	0.00%
mmap	21	0.118	0.003	0.006	0.009	5.66%
mprotect	4	0.030	0.004	0.008	0.012	23.08%
munmap	1	0.010	0.010	0.010	0.010	0.00%
brk	3	0.009	0.002	0.003	0.004	23.94%
rt_sigaction	3	0.005	0.001	0.002	0.002	14.85%
access	1	0.006	0.006	0.006	0.006	0.00%
Dup2	2	0.005	0.002	0.002	0.003	8.43%
arch_prctl	1	0.002	0.002	0.002	0.002	0.00%



# Perf probe

- Dynamically created tracepoints

```
# perf probe _copy_from_user
Added new event:
  probe:_copy_from_user (on _copy_from_user)

You can now use it in all perf tools, such as:

    perf record -e probe:_copy_from_user -aR sleep 1
```

- Count or sample

```
# perf stat -a -e probe:_copy_from_user sleep 1

Performance counter stats for 'system wide':

      1,209          probe:_copy_from_user

1.001884191 seconds time elapsed
```



# Scripting

- `perf script`
- `perf script -l`
- `perf script -g [python|perl]`
- `perf script record <script> [perf record options]`
- `perf script report <script> [script options]`







# Advanced topics


# Flame Graphs

<http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html>

- `perf record -a -g -e cycles:k scp /tmp/file.10gb host:`
- `perf script | stackcollapse-perf.pl | flamegraph.pl > data.svg`
- `firefox data.svg`



# Flame Graphs



# PAPI / self-monitoring

- Performance Application Programming Interface
- <http://icl.cs.utk.edu/papi/>
  - Has it's own event monikers termed “event presets”s but can work with perf raw event codes.
- Alternative: see implementation of `tools/perf/tests/rdpmc.c`
  - [http://web.eece.maine.edu/~vweaver/projects/perf\\_events/overhead/2015\\_ispass\\_overhead.pdf](http://web.eece.maine.edu/~vweaver/projects/perf_events/overhead/2015_ispass_overhead.pdf)
  - Mainline commits `fe4a3308` and `08aa0d1f`

# Off-cpu analysis

- Monitoring time spent not running.
- Many reasons
  - disc io
  - synchronization (locks, mutexes, signals)
  - virtual memory
  - context switching
- <http://www.brendangregg.com/offcpuanalysis.html>



# Virtualization

- Monitoring guest requires virtualization of pmu hardware.
- Xen DomU reports software/tracepoint only
- QEMU: use '-cpu \*,pmu=on' or '-cpu host' (kvm)
- Use perf kvm to monitor guest






Questions?

Thanks for listening. Further questions welcomed via email. Bug reports are great too.

**Thank you.**







**Have a Lot of Fun, and Join Us At:**

**[www.opensuse.org](http://www.opensuse.org)**

